

SPEC vs Agentic, one vote per workload — combined agent fence (tool + harness)

SPEC = 26 benchmarks, Agentic = the revised all-resolved 36; each workload contributes its per-window MEDIAN once (agent windows: the two fences' raw counts summed) · label = group median · black bar = median, white diamond = mean

