

suite benchmark (toy dataset)


● same software on a realistic dataset

--- SPEC median (26)


--- Agentic median (36)

one value per workload = metric over its whole runtime


IPC




Branch MPKI




Branch-direction MPKI




BTB MPKI (BAClears)




L1I MPKI (code-read)




uop-cache (DSB) MPKI




DSB coverage (%)




L1D-load MPKI




L2-load MPKI



LLC MPKI



DRAM read (GB/s)



Context switches (/CPU-s)

