



suite benchmark (toy dataset)





same software on a realistic dataset




SPEC median (26)

Agentic median (36)


one value per workload = metric over its whole runtime




L1I MPKI (code-read)




uop-cache (DSB) MPKI




DSB coverage (%)




L1D-load MPKI




L2-load MPKI



LLC MPKI



DRAM read (GB/s)



Context switches (/CPU-s)

